

Rohan Verma

Machine Learning Engineer

Hyderabad, India | rohan.verma.mle@gmail.com | +91 90123 45678 | linkedin.com/in/rohan-verma-mle | github.com/rverma-ml

PROFESSIONAL SUMMARY

Machine Learning Engineer with 5 years of experience building and productionizing ML systems at scale, including recommendation engines, NLP pipelines, and computer vision models. Strong foundation in software engineering and MLOps, with a focus on model reliability, latency optimization, and reproducible training pipelines in cloud-native environments.

TECHNICAL SKILLS

Languages: Python, Go, C++, SQL

ML Frameworks: PyTorch, TensorFlow, ONNX, Scikit-learn

NLP/CV: Hugging Face Transformers, spaCy, OpenCV, YOLO

MLOps: Kubeflow, MLflow, Docker, Kubernetes, Airflow, Feast (Feature Store)

Cloud & Infra: AWS (EC2, S3, SageMaker, Lambda), GCP (Vertex AI), Terraform

Serving & APIs: FastAPI, gRPC, Triton Inference Server, Redis

PROFESSIONAL EXPERIENCE

Machine Learning Engineer II — Swiggy

Mar 2022 – Present

Hyderabad, India

- Designed and shipped a restaurant recommendation model (two-tower neural network) serving 25M+ daily active users with p99 latency under 80ms.
- Built a real-time feature store using Feast and Redis, cutting feature computation duplication across 6 ML teams by 60%.
- Optimized model inference using ONNX Runtime and quantization, reducing GPU serving costs by 35%.
- Implemented automated retraining pipelines with Kubeflow, reducing model refresh cycle from weekly manual runs to nightly automated runs.
- Introduced canary deployment and shadow-testing framework for new models, catching 3 production regressions before full rollout.

Machine Learning Engineer — Freshworks

Aug 2019 – Feb 2022

Chennai, India

- Built an NLP-based ticket classification and routing system (fine-tuned BERT) that reduced average first-response time by 27%.
- Developed a document extraction pipeline using OCR and named entity recognition for customer support automation.
- Containerized and deployed ML services with Docker and Kubernetes on AWS EKS, achieving 99.9% uptime.
- Set up CI/CD pipelines for model versioning and automated testing using GitHub Actions and MLflow.

KEY PROJECTS

Open-Source Contribution — Hugging Face Transformers | Python, PyTorch

- Contributed a bug fix and performance improvement to the tokenizers library, merged into the main branch.

EDUCATION

B.Tech in Computer Science and Engineering*2015 – 2019*

Indian Institute of Information Technology (IIIT), Hyderabad

*CGPA: 8.7/10***CERTIFICATIONS**

- TensorFlow Developer Certificate (2021)
- Google Cloud Professional Machine Learning Engineer (2023)